


# REPRODUCIBLE CONTAINERS FOR SCIENTIFIC COMPUTING

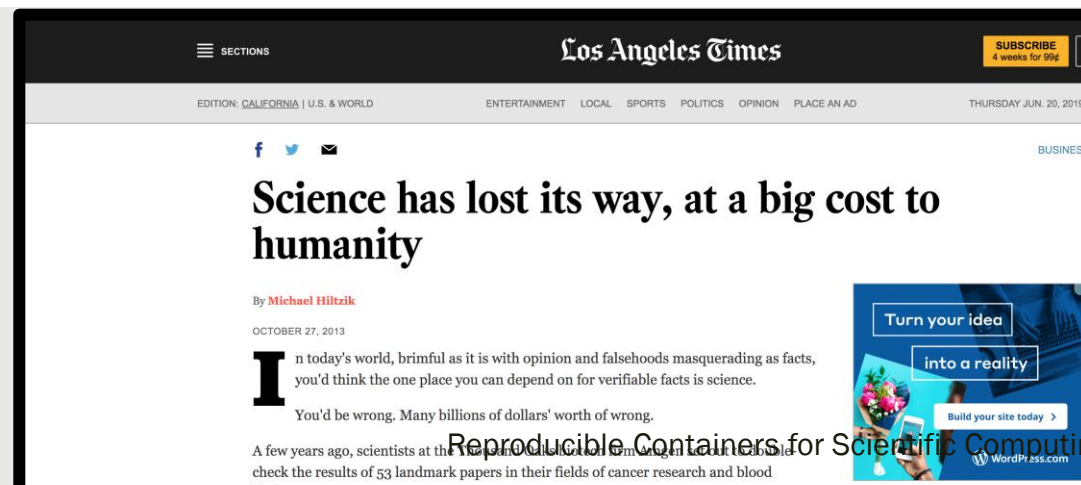
Tanu Malik

The DICE Laboratory  
*School of Computing*  
*DePaul University, Chicago IL*







# Reproducibility and Replicability in Sciences: A growing concern



# Concerns across sciences



Source: 1500 scientists lift the lid on reproducibility. *Nature Survey*. Accessed 25th May, 2016





# Concerns in Computer Sciences

- Definition: Given an experiment described at  $t$ , Can it be downloaded at  $t'$ , and its source code be built within a reasonable amount of time.

N = 402 experiments

Response	Percentage	# of experiments
No response to code requests	36.1%	145
Declined to provide code	10.2%	41

N = 219 experiments

Response	Percentage	# of experiments
Failed to build	5.02%	11
Built $\leq$ 30 minutes	32.4%	71
Built $>$ 30 minutes	15.9%	35
Reasonable effort	46.6%	102

T. Proebsting, A. M. Warren, and C. Collberg. 2015.

Repeatability and benefaction in computer systems research. University of Arizona TR 14. Vol. 4. 1–68.





# Challenges in Reproducibility and Replicability

- **Reproducibility:** Obtaining consistent results when using the same or similar input data, computational steps, conditions of analysis, etc.
- **Replicability:** Obtaining consistent results when using different input data, computational steps, conditions of analysis, etc.

## Challenges

- Need for **guarantees-based or statistical-based methods** for conduct of reproducible research
- Need for **infrastructure** that supports reproducible research
- Need for **policies** that incentivize and enforce reproducible research

National Academies of Sciences, Engineering, and Medicine 2019. Reproducibility and Replicability in Science. Washington, DC: The National Academies Press. <https://doi.org/10.17226/25303>.





# Foundations of Reproducible Scientific Computing

## Guarantees

*Lineage-based methods*

*Replay:* VLDB'22

*Debugging:* TaPP'20, HiPC'22

*Guarantees:* TaPP'13, ICDE'15, MDPI'18

## Infrastructure

*Container-based Tools*

*w/ Lineage:* eScience'17, eScience '19, eScience'22

*Size Reduction:* HiPC'20, Access'23,

*Vs Docker:* ICCS'15

*Documenting:* PARCO'20

## Policies

*Artifact Evaluation*

*Surveys:* PRECS'22, IEEE CiSE'21





# Foundations of Reproducible Scientific Computing

## Guarantees

### *Lineage-based methods*

**Replay:** VLDB'22

**Debugging:** TaPP'20, HiPC'22

**Guarantees:** TaPP'13, ICDE'15, MDPI'18

## Infrastructure

### *Container-based Tools*

**w/ Lineage:** eScience'17, eScience '19, eScience'22

**Size Reduction:** PARCO'20,

HiPC'20, Access'23,


**Vs Docker:** ICCS'15

**Documentation:** PARCO'20

## Policies

### *Artifact Evaluation*

**Surveys:** PRECS'22, IEEE CiSE'21





# Share and Reproduce an Application




Alice



Bob

Alice wants to share her input data files and program source code with Bob  
Bob wants to reproduce Alice's application to validate her outputs.





# Repeat Share and ~~Reproduce~~ an Application




Alice



Bob


Alice wants to share her input data files and program source code with Bob  
Bob wants to repeat Alice's application to validate her inputs and outputs.





# Alice's sharing options


1. Email a tar/gzip
2. Build a website with model code, parameters, and data
3. Create a virtual machine or container





# Not sufficient for reproducible research

- ~~1. Email a tar/gzip~~
- ~~2. Build a website with model code, parameters, and data~~
- ~~3. Create a virtual machine or container~~






# Not sufficient for reproducible research

- ~~1. Email a tar/gzip~~
- ~~2. Build a website with model code, parameters, and data~~

Missing environment files: No isolation guarantee

- ~~3. Create a virtual machine or container~~

Failing rebuilds: No repeatable guarantee





# Docker for repeating an application?

```
# syntax=docker/dockerfile:1
FROM golang:1.16-alpine AS build

# Install tools required for project
# Run `docker build --no-cache .` to update dependencies
RUN apk add --no-cache git
RUN go get github.com/golang/dep/cmd/dep


# List project dependencies with Gopkg.toml and Gopkg.lock
# These layers are only re-built when Gopkg files are updated
COPY Gopkg.lock Gopkg.toml /go/src/project/
WORKDIR /go/src/project/
# Install library dependencies
RUN dep ensure -vendor-only

# Copy the entire project and build it
COPY . /go/src/project/
RUN go build -o /bin/project

# This results in a single layer image
FROM scratch
COPY --from=build /bin/project /bin/project
ENTRYPOINT ["/bin/project"]
CMD ["--help"]
```

Because this line installs the most recent version upon rebuilding


This line no longer rebuilds





# Can we compose Containerization and Lineage?

<http://sciunit.run>






# Sciunit: Compose Containerization and Lineage

<http://sciunit.run>

*Key Idea:* Identify and isolate data dependencies during program execution and infer lineage between dependencies

D.H. Ton That, G. Fils, Z. Yuan, T. Malik. Sciunits: Reusable Research Objects.  
In IEEE eScience Conference (eScience), 374-383, 2017


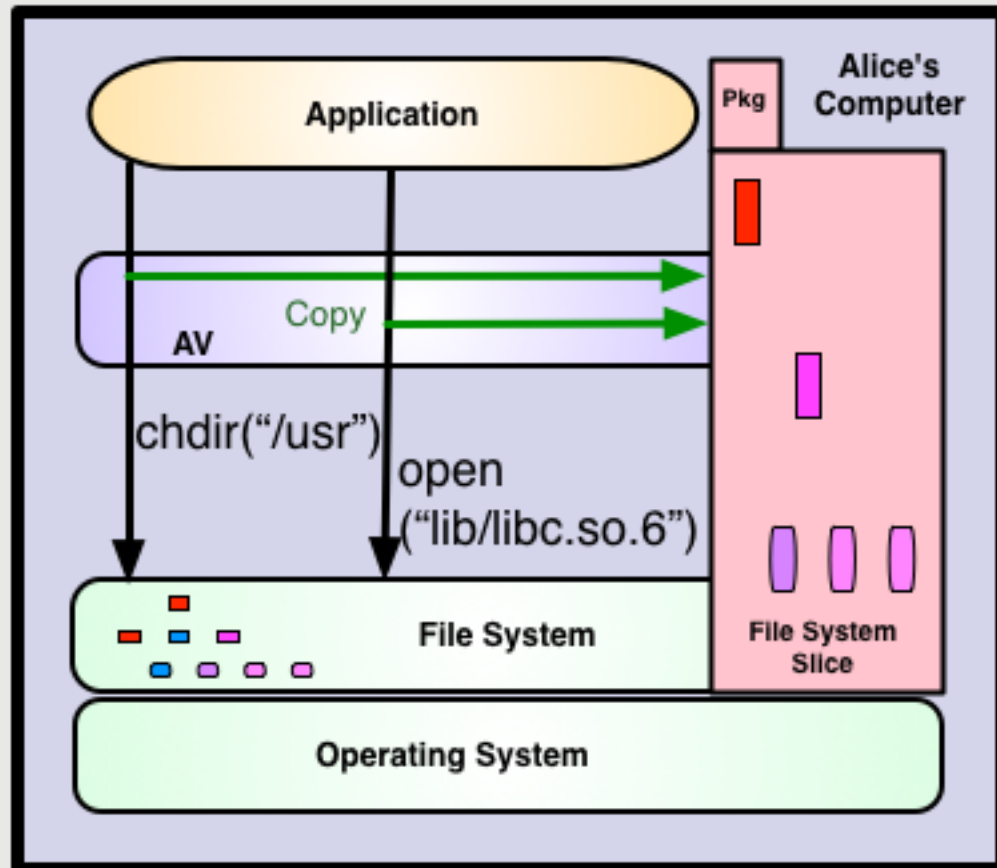




# Host: Use *ptrace* to observe executions

## Audit Phase

- Audit ~50 system calls related to process control, file I/O, and network
  - If file is /dev/random capture return bytes as well
- At the time of interception:
  - Generate an execution trace of system call events in real-time
  - Copy files mentioned as part of system calls into a container

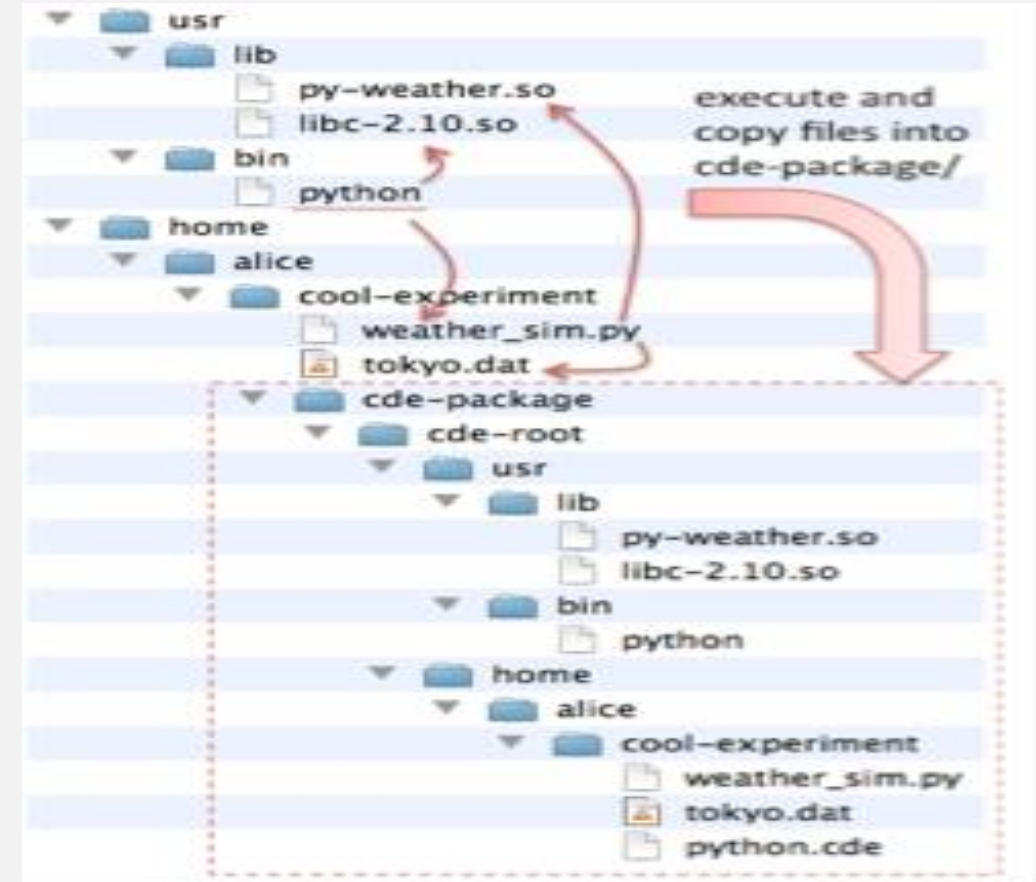
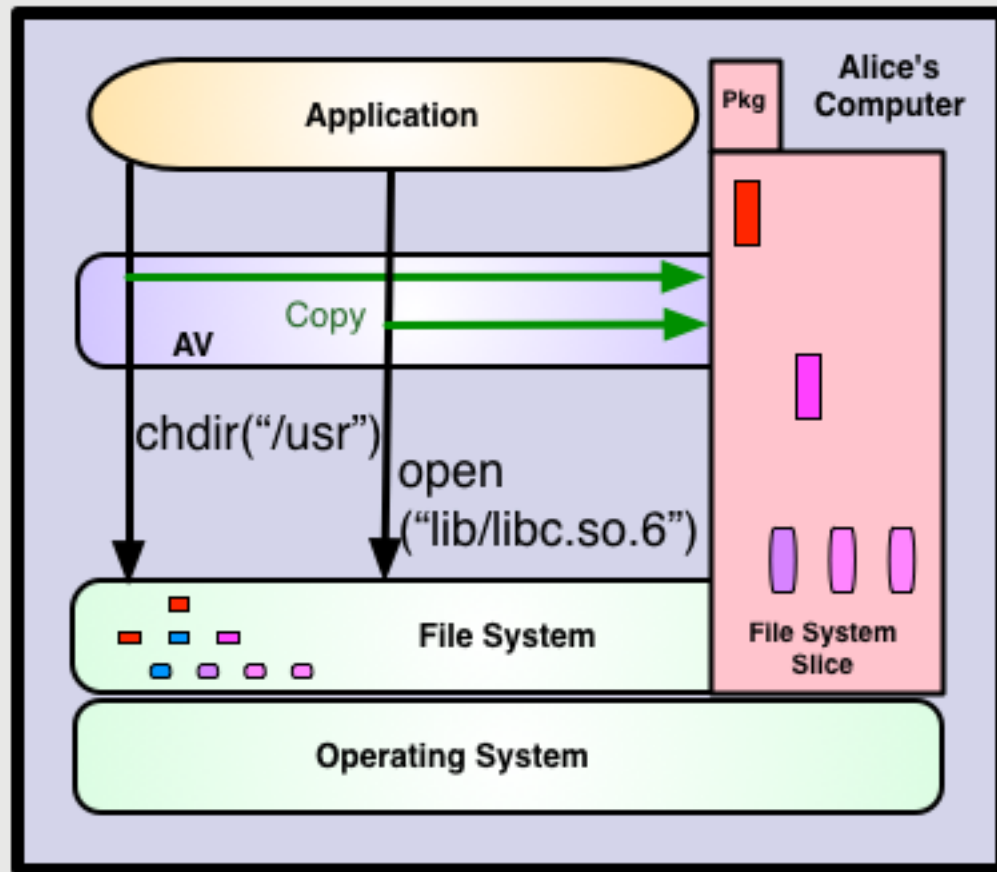






# Create a *chroot*-based container

- Audit provenance during container creation time







# Alices shares *sciunits* and Bob repeats them




sciunits






# Target: Use *ptrace* to redirect executions into the container





## Repeat Phase


- Redirection during repetition is only for file- and network-related system events.
- Repeat execution happens within a process and file namespace.



# Mapping system events to a provenance graph

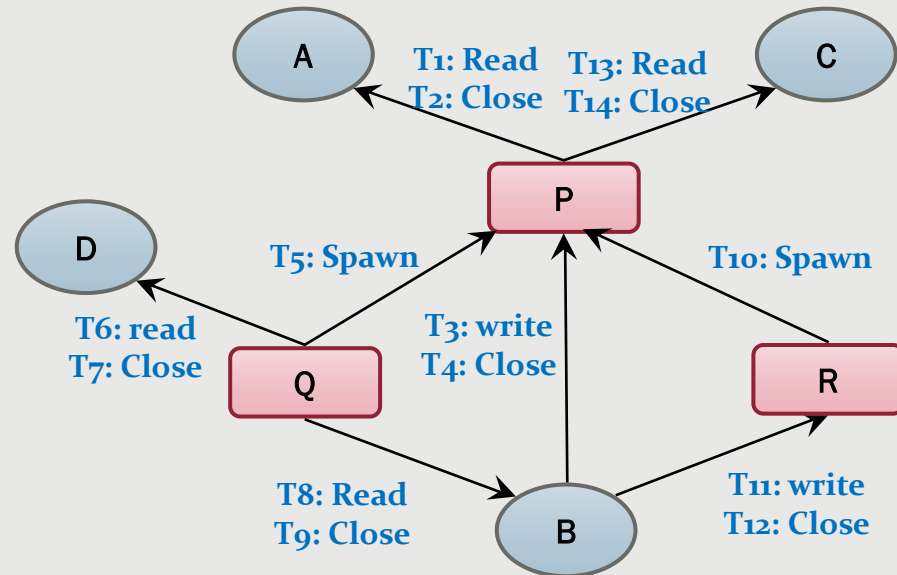


Z. Yuan, D.H. Ton That, S. Kothari, G. Fils, T. Malik. Utilizing Provenance in Reusable Research Objects, In *Special Issue on Using Computational Provenance*, MDPI Informatics, Vol 5(1), 2018.


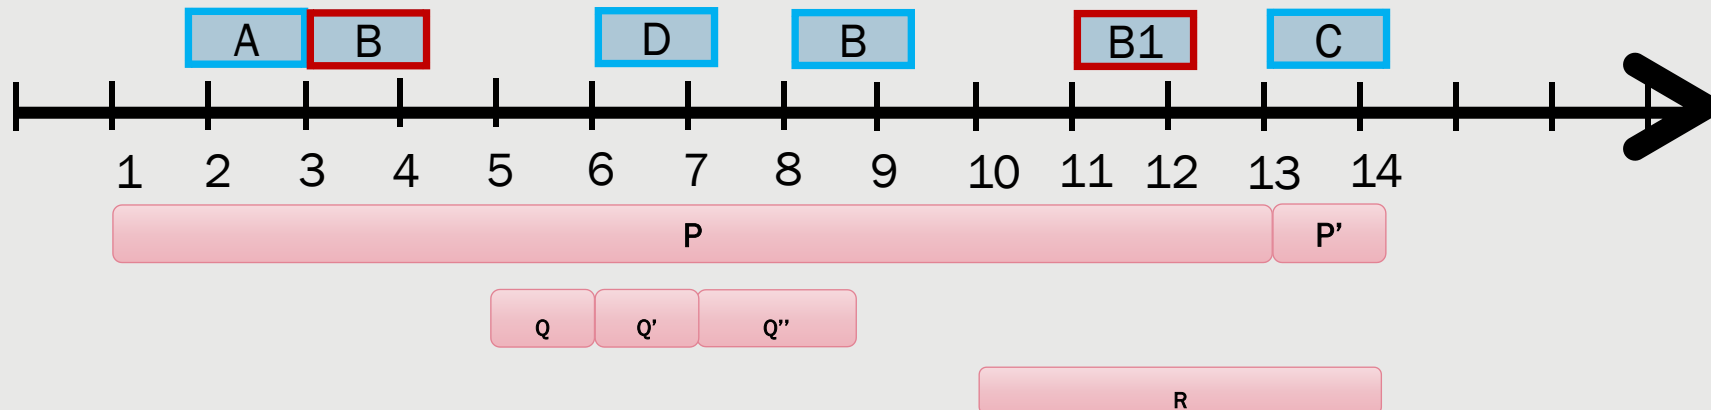




# Versioning of process and file nodes

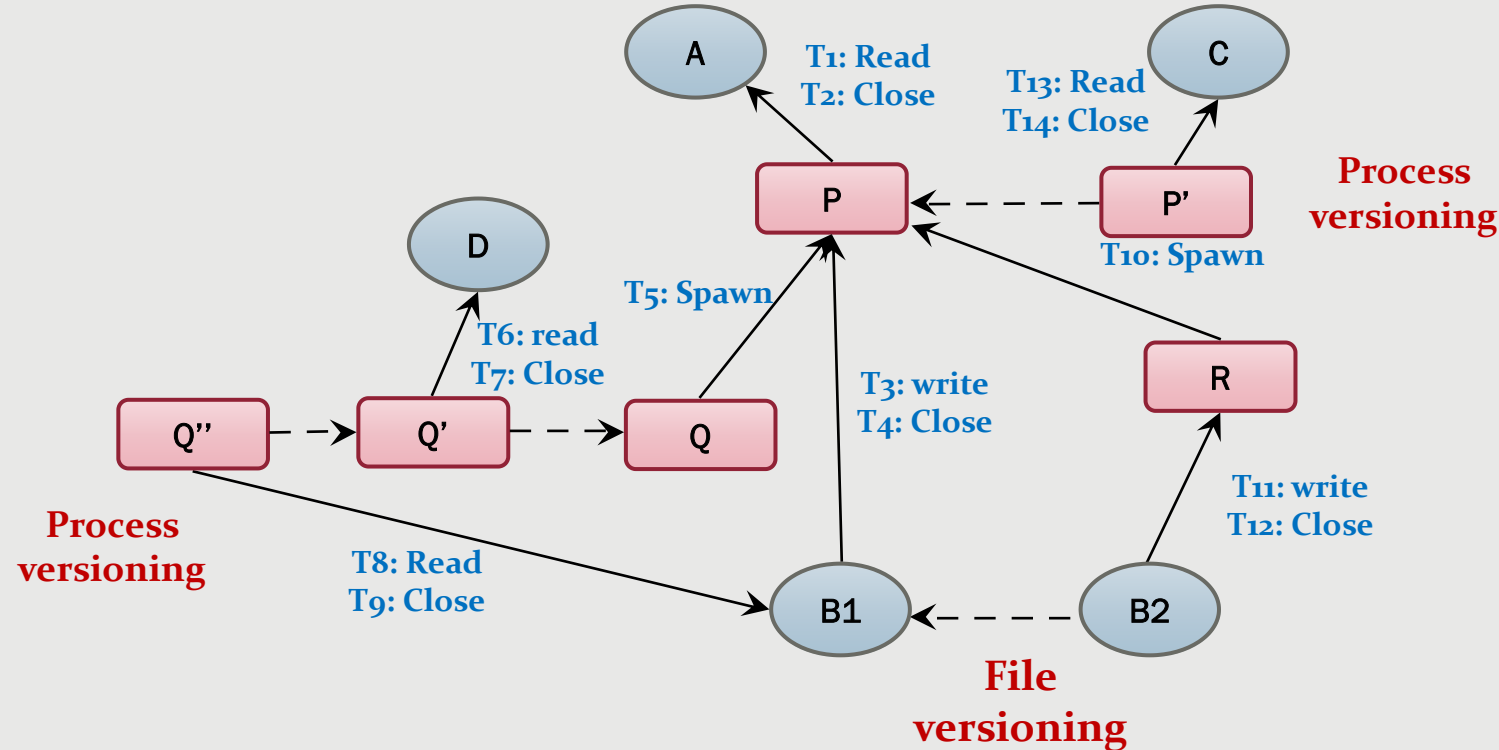


- Each file is versioned if it is written over
  - Write on a file changes file content
- Each process if versioned if it reads a new file
  - Read in a process changes process state






# Inferencing causality of nodes




- If a file is read later, it is not causal on files written before.
  - B1 or B2 is not causal on C
- Changes to files can determine the causal graph that is impacted.
  - Changing A will impact all nodes, but the effect of C and D





# Execution Trace

- Let  $P$  be a program
- A execution trace  $L$  for  $P$  is a 2-tuple  $\langle G, R \rangle$
- Provenance  $G = (V, E, T)$  with nodes  $V$  and edges  $E \subseteq V \times V$ .
  - Each node  $v \in V$  and edge  $e \in E$  has annotations
  - $T : E \rightarrow T \times T$  is a function mapping edges to intervals from a discrete time domain  $T$
- Package  $(R, <)$  in which elements  $R \subseteq V$  are organized as a tree s.t  $r_i \in R$  maps to a content in  $v \in V$  and  $v$  has an outgoing edge






# L is a repeatable iff

- L is deterministic, i.e., it will lead to the same result,  $R = \{R_1, \dots, R_n\}$ , every time  $t' > t$
- $L(G)$  includes necessary activities, entities and edges,
  - *it does not leave out activities, entities and edges in  $G(V,E,T)$  that may have caused  $R$*
- $L(R)$  is sufficient
  - *does not include superficial elements that do not cause  $R$*

Q. Pham, T. Malik, B. Glavic, I. Foster. Light-weight Database Virtualization. In *IEEE International Conference on Data Engineering, ICDE*, 2015.





# Evaluation



- Use cases:
  - FIE: Chicago Food Inspections Evaluation (~ 307 MB)
    - *A ML prediction model of food inspections*
  - VIC: Variable Infiltration Capacity (~ 1.2 GB)
    - *VIC: A Hydrology application*
  - IQE: Incremental Query Execution (~ 22 MB)
    - *A DB application with incremental query processing*
- Base Lines:
  - Docker
  - IncPy





# Real Applications--Description


- FIE: A ML prediction model of food inspections
- VIC: A Hydrology application
- IQE: A DB application with incremental query processing

TABLE I: Usecases descriptions.


	FIE	VIC	IQE
Source code languages	R, Bash	C, C++, Python, C shell script, Fortran	Python
Source code files	29	97	5
Data files	14	11,481	5
Dependency files	659	357	112
Size of all files	306.6 MB	1.2 GB	22 MB
Normal run time	286.756 s	40.259 s	5.226 s




# Container size




1. Docker containers are 19X, 7X and 2.5X larger than those of Sciunit.
2. Sciunit containers are even slightly smaller than the original application package size.



# Auditing and re-execution time



1. Docker spends 4X and 9X longer time than Sciunit to build FIE and VIC container, 7X and 2.5X larger.
2. Sciunit only spends slightly higher time than the original execution time to build containers.






# Sample Interaction

```
1. > sciunit create FIE
2. > sciunit exec ./FIE.sh ./DATA/weather_201710.Rds
    0. Download...
    1. Calculate violation matrix...
    2. Calculate heat map...
    3. Generate model data with ./DATA/weather_201710.Rds...
    4. Apply random forest model...
    5. Evaluation...
3. > sciunit list
    e1 Dec 4 12:44 ./FIE.sh ./DATA/weather_201710.Rds
4. > sciunit show
    id: e1
    sciunit: FIE
    command: ./FIE.sh ./DATA/weather_201710.Rds
    size: 306.6 MB
    started: 2017-12-04 12:44
5. > sciunit push
    ...
    Title for the new article: FIE
    new: 306.6 MB [01:05, 4.72MB/s]
6. > sciunit copy
    mSLLTj#
```

Alice's Computer

```
1. > sciunit repeat e1
    ...
    0. Download...
    1. Calculate violation matrix...
    2. Calculate heat map...
    3. Generate model data with ./DATA/weather_201710.Rds...
    4. Apply random forest model...
    5. Evaluation...
    ...
2. > sciunit repeat e1 <27050>
    ...
    3. Generate model data with ./DATA/weather_201710.Rds...
    ...
3. > sciunit given '/tmp/weather_201801.Rds' e1 %
    ...
    0. Download...
    1. Calculate violation matrix...
    2. Calculate heat map...
    3. Generate model data with /tmp/weather_201801.Rds...
    4. Apply random forest model...
    5. Evaluation...
    ...
```


Bob's Computer





# Summary

- FLINC: Notebook-based containerization
  - CHEX: Notebook-based multiversion replay
  - Kondo: Reducing the size of data-intensive containers.
  - ProvScope: Container-based differencing
  - Sched: Scheduling of application-virtualized containers.
- 
- Please check out our work at:
    - <https://dice.cs.depaul.edu/>
    - <https://github.com/depaul-dice>





# Thank You

- Please check out our work at:
  - <https://dice.cs.depaul.edu/>
  - <https://github.com/depaul-dice>
- Questions/Contact: [tanu.malik@depaul.edu](mailto:tanu.malik@depaul.edu)

